

Methods

2021-12-13

Methods

General description

We built an R (R Core Team 2020) based framework to simplify two aspects of systematic reviews: record acquisition and classification. The code used to generate the results in the manuscript is available at X while an update version of the framework is distributed as an R package at Y. The framework includes several modules that communicate through intermediate outputs stored in standard formats, which make it possible for users to extend the framework or easily integrate it with other tools in their pipeline. See Supplemental Material S1 for an in-depth description of the framework and how to use it.

The tasks carried out by the framework are grouped into “sessions,” which comprise obtaining scientific citation data (records) using a search query and then labelling them as relevant (“positive” in the rest of the text) or not (“negative”) for the topic of interest with the help of a machine learning engine (Fig. 1). The initial search query should be built using domain knowledge, trying to achieve a high relevant/non-relevant record ratio.

The framework can then generate a new data-driven query from this labelled set to perform a new session to find records possibly missed by the first query.

Record’s acquisition and initial labelling

We built a set of tools to let users automatically search and download citation data from three major scientific databases (“sources”): Pubmed (<https://pubmed.ncbi.nlm.nih.gov/>), Web Of Science (WOS, <https://apps.webofknowledge.com/>) and the Institute of Electrical and Electronics Engineers (IEEE, <https://ieeexplore.ieee.org/Xplore/home.jsp>). The framework takes care of authorization management for non-open databases like WOS and IEEE. It is also possible to download and import records in the framework manually; this is particularly useful to acquire records from the SCOPUS (<https://www.scopus.com/>) and EMBASE databases (<https://www.embase.com/>), for which a comprehensive API interface was not easy to build. An extra manual search was necessary also for Pubmed, since the API and the web interface have different rule expansion algorithms and return slightly different results (“NCBI Insights : Updated Pubmed e-Utilities Coming in April 2022!” n.d.). A short guide on how to set up the framework for each supported database is available in Supplemental Material S3.

The acquired records are merged into a single database, resolving duplicates and different formatting between sources. The records are ordered according to the frequency of the positive query terms (e.g., not preceded by a *NOT* modifier) in the title and abstract (“simple query ordering”).

The researcher is then asked to label a number of records to create the “initial training set” needed to start the automatic classification. We suggest manually labelling the first 250 records (see “hyperparameter optimization” later). The simple query ordering increases the positivity rate in the initial training set (Wallace et al. 2010), which provide higher sensitivity during automatic classification (Chawla, Japkowicz, and Kotcz 2004).

Text feature extraction

The collected citation data have a number of fields characterising a scientific publication. The framework models the relevance of a record based on the following fields: title, abstract, authors, keywords, MESH terms

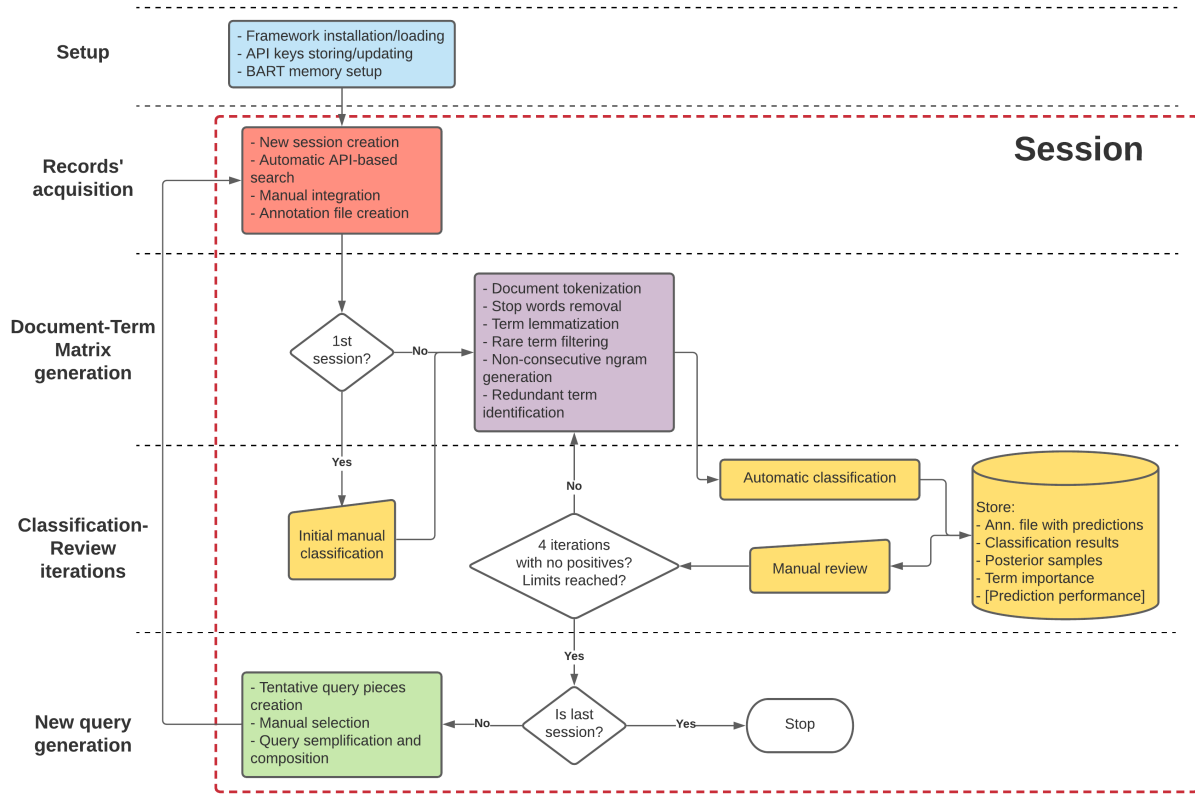


Figure 1. Framework's visual depiction.

(Lipscomb 2000). A series of Natural Language Processing (NLP) techniques (Baeza-Yates, Ribeiro-Neto, and others 1999; Marshall and Wallace 2019; Ananiadou and McNaught 2006) are employed to transform the textual information in these fields into features for machine learning through a bag-of-words approach (Marshall and Wallace 2019). The processing of free text fields (title, abstract) includes: tokenization (i.e., extracting the terms), common stopwords (i.e., sentence components bringing no meaning) removal, part-of-speech filtering (only nouns, adjectives, verbs and untagged terms are retained), and lemmatization of the terms (i.e., reduction to their base grammar form). Text processing for authors, keywords and MESH terms identifies logical units (e.g., author’s full names, composite keywords) and extracts them. Terms appearing in less than 5% of the labelled documents are removed from negative records. All terms in the positive set are kept to increase sensitivity at the cost of specificity. Some terms tend to co-appear in records (non-consecutive ngrams, nc-ngrams), often carrying a particular meaning when copresent. To detect nc-ngrams, we generated a word network representation (Rousseau 2015) posing edges between terms with a cosine similarity in terms of record co-occurrence > 0.5 . We extracted the maximal cliques in the network (Eppstein, Löffler, and Strash 2010) representing highly correlated groups of terms; These groups of terms are added to the data set as individual features. Only nc-ngrams comprising a maximum of ten terms are kept. A second network is built using a co-occurrence threshold of 0.9. In this case, the cliques represent terms that always appear together and therefore can be considered redundant (i.e., they do not need to be considered separately). These terms are merged to increase computation efficiency and reduce overfitting. The output is a Document-Term Matrix (DTM), with N_d rows representing the records (D_i), N_t terms column for the t_{field} terms (divided by record field) and 0, 1 values whether $t_{field} \in D_i$. We also enriched the DTM with features referencing the number of terms in each field to help the model scale term importance based on the field length.

Label prediction

We used a Bayesian Additive Regression Trees (BART) machine learning “classification model” (Chipman et al. 2010) (in the implementation of Kapelner and Bleich 2013) to predict the probability of a record of being relevant, given the information coded into the enriched DTM and the initial training set. We set up the BART model to use 2000 MCMC iterations (after 250 burn-in iterations) and 50 trees; we used a k value of 2 to regularize extreme prediction and let the model use missing fields in the DTM as features (Kapelner and Bleich 2015). Positive records are oversampled ten times to increase sensitivity (Batista, Prati, and Monard 2004).

The output is a posterior predictive distribution (PPD) for each record describing its probability of being relevant (i.e., a positive match). An ensemble of ten models was fitted to improve prediction stability by averaging the PPD between models (Zhou 2021; Dietterich 2000).

To assign the labels, we employed an “active learning” Miwa et al. (2014) approach, where a human reviews a specific subset of predictions made from the machine, which is then retrained on the manually reviewed data set. This process proceeds iteratively, decreasing prediction uncertainty.

Label assignment is done through the identification of an “uncertainty zone” whose construction is possible thanks to the Bayesian nature of BART, which provides full PPDs instead of point-wise predictions for each record.

To describe the process formally, we define

$$\pi_i = \frac{1}{M} \sum_{j=1}^M Pr(L_i = 1 | DTM, m_j)$$

as the PPD of a record D_i being relevant (i.e, having a positive label, $L_i = 1$), averaging the PPDs of the ensemble of $M = 10$ models m , and

$$\begin{aligned}\pi_{i,l} &= \{\pi_i : Pr(\pi_i) = 1\%\} \\ \pi_{i,u} &= \{\pi_i : Pr(\pi_i) = 99\%\}\end{aligned}$$

as respectively the lower and upper boundaries of the 98% quantile interval of π_i (98% Predictive Interval, 98% PrI).

Then we identify the “uncertainty zone” as

$$U_\pi = [\max \vec{\pi}_u^-, \min \vec{\pi}_l^+]$$

with $\vec{\pi}_u^-$ being the vector of $\pi_{i,u}$ with a negative label and $\vec{\pi}_l^+$ the vector of $\pi_{i,l}$ with a positive label. That is, U_π defines a range of values between the smallest $\pi_{i,l}$ in the set of already labelled positive records L_p and the largest $\pi_{i,u}$ related to the negative ones L_n , noting that the two limits can appear in any order.

Consequently, a record D_i will be labelled as positive if

$$\pi_{i,l} > \max_{\pi \in U_\pi} \pi$$

that is, the record lower 98% PrI boundary should be higher than every value in the uncertainty zone. In other words, for a record to be labelled positive, its PPD should be within the range of the mixture of PPD of the previously labelled positive records and not cross the distributions of the negative records.

Conversely, a record is labelled as negative if

$$\pi_{i,u} < \min_{\pi \in U_\pi} \pi$$

All other records are labelled as “uncertain.”

Manual review is then necessary for: 1) uncertain records, 2) positive records (to avoid false positives), 3) records whose predicted label differs from the existing manual one. The last case helps identify human errors or inconsistent labelling criteria.

The automatic classification and the manual review steps alternate in a loop (CR iterations) until no new positive matches are found in four consecutive iterations.

Relevant term extraction

BART model it is possible to extract the proportion of times a term was used in a posterior tree over the sum of total inclusions of all variables (Kapelner and Bleich 2013), as a measure of feature importance for the model. We extracted the terms, the part of the citation data they were used in, the average “inclusion rate” among the ensemble models (over 10,000 inclusions) and its ratio over the standard deviation of this inclusion (inclusion stability, IS). For each term we ran a Poisson regression to get its linear association with a positive label and reported it as Relative Risk (RR) with the number of standard errors as significance index (Statistic); the comparison between the inclusion rate in the BART models and the linear association allows to spot relevant non-linear effects (i.e, the feature is relevant only in association with others). We reported in the Results only the first (in order of inclusion rate) 15 terms with $IS > 1.5$ while the first fifty terms regardless of inclusion stability are reported in Supplemental Material S2.

New search query generation

We created an algorithm that generates a new search query to acquire further relevant publications missed during the first search, possibly at a reasonable cost in specificity (i.e., a higher number of negative results). The algorithm encompasses a number of steps:

- We fit a partition tree (Therneau and Atkinson 2019) between the DTM and 800 samples from the PPD; if a term is present multiple times in the DTM (e.g., both title and abstract), they are counted just one, and field term count features are removed. This step generates a list of rules composed by *AND/NOT* “conditions” made of terms/authors/keywords/MESH tokens, which together identify a group of records.
- For each rule, negative conditions (i.e., *NOT* statements) are added iteratively, starting from the most specific one, until no conditions are found that would not also remove positive records.
- The extended set of rules is sorted by positive-negative record difference in descending order. The cumulative number of unique positive records is computed and used to group the rules. Rules inside each group are ordered by specificity.
- The researcher is then asked to review the rule groups, selecting one or more rules (useful if they convey different meaning) from each, or edit them (in case too specific positive or negative conditions were included). It is possible to exclude a group of rules altogether, especially those with the worse sensitivity/specificity ratio.
- The selected rules are joined together by *OR* statements, defining a subset of records with a sensibly higher proportion of positive records than the original one.
- Redundant rules (i.e., rules whose positive records are already included in more specific ones) and conditions (i.e., conditions that once removed do not decrease the total number of positive or do not increase the negative records) are removed.
- Finally, the rules are re-elaborated in a query usable on the major scientific databases.

Since the algorithm is data-driven, it creates queries that effectively select positive records from the input data set but may be not specific enough once applied to actual research databases. Therefore we appended an extra subquery in *AND*, which specifies the general topics of our search and delimitates the search domain. The new query was used to initiate a second search session.

Performance evaluation

We trained a simple Bayesian logistic regression (surrogate model) on the reviewed records to evaluate the classification model consistency (see Discussion for the theoretical justification). The surrogate model uses as predictor the lower bound of the 98% PrI of the records’ PPD with weakly regularizing, robust priors for the intercept (Student T with $\nu = 3, \mu = 0, \sigma = 2.5$) and the linear coefficient (Student T with $\nu = 3, \mu = 0, \sigma = 1.5$).

The quality of the model was evaluated through Bayesian R^2 (Gelman et al. 2019), of which we reported the posterior median and 90% Credible Interval [90% CrI]. The R^2 also provides an evaluation of the consistency of the original classification model. Given that this model is conditional only on the BART predictions and not on the DTM, it is characterized by more uncertainty, providing plausible worst-case scenarios.

The surrogate model is then used to generate the predictive cumulative distribution of the number of total positive records in the whole data set. This distribution allows estimating the expected total posterior “Sensitivity” and “Efficiency” of the classification model in the full (unreviewed) data set. Efficiency is summarised by the “Work saved over random” (WSorR) statistic: one minus the ratio between the number of record manually reviewed and those that would be required to find the same number of positives if classification was performed choosing records randomly; this last quantity is estimated through a negative hypergeometric distribution (Chae 1993) over the predicted amount of positive records.

For the number of predicted positive records, the sensitivity and the efficiency, we reported the “truncated 90% PrI” [trunc. 90% PrI], which is the uncertainty interval bounded at the number of observed total positive records (i.e., there cannot be less predicted positive records than observed).

Hyperparameter evaluation

Our classification algorithm has a number of hyperparameters:

- Size of the initial training set: 50, 100, 250, 500 records;
- Number of models in the ensemble: 1, 5, 10, 20, 40, 60 repetitions;
- Oversampling rate of positive records: (1x (i.e., no oversampling), 10x, 20x;

- PrI quantiles for building the uncertainty zone: 80%, 90%, 98%;
- Source of randomness between models in the ensemble: MCMC sampling only (Robert, Casella, and Casella 2004), MCMC plus data bootstrapping (Breiman 1996) of the training set.

To evaluate the hyperparameter effect of performance, we set up a “grid search” (Claesen and De Moor 2015; Yang and Shami 2020) on a prelabelled “validation set” made of the first 1200 records from the first session data set. The framework tested each hyperparameter combination until four CR iterations with no positive records were returned or the whole data set got labelled.

For each combination, a performance score was computed as the product of “Efficiency” (1 minus the ratio of records that required review over the total) and “Sensitivity” (number of positive records found over the total of positives). We then identified homogeneous “performance clusters” of scores given hyperparameter values using a partition tree (Therneau and Atkinson 2019). For the rest of the study, we used the best hyperparameter set in terms of Sensitivity followed by Efficiency from the cluster with the highest average score.

- Ananiadou, Sophia, and John McNaught. 2006. *Text Mining for Biology and Biomedicine*. Citeseer.
- Baeza-Yates, Ricardo, Berthier Ribeiro-Neto, and others. 1999. *Modern Information Retrieval*. Vol. 463. ACM press New York.
- Batista, Gustavo EAPA, Ronaldo C Prati, and Maria Carolina Monard. 2004. “A Study of the Behavior of Several Methods for Balancing Machine Learning Training Data.” *ACM SIGKDD Explorations Newsletter* 6 (1): 20–29.
- Breiman, Leo. 1996. “Bagging Predictors.” *Machine Learning* 24 (2): 123–40.
- Chae, Kyung-Chul. 1993. “Presenting the Negative Hypergeometric Distribution to the Introductory Statistics Courses.” *International Journal of Mathematical Education in Science and Technology* 24 (4): 523–26.
- Chawla, Nitesh V, Nathalie Japkowicz, and Aleksander Kotcz. 2004. “Special Issue on Learning from Imbalanced Data Sets.” *ACM SIGKDD Explorations Newsletter* 6 (1): 1–6.
- Chipman, Hugh A, Edward I George, Robert E McCulloch, and others. 2010. “BART: Bayesian Additive Regression Trees.” *The Annals of Applied Statistics* 4 (1): 266–98.
- Claesen, Marc, and Bart De Moor. 2015. “Hyperparameter Search in Machine Learning.” *arXiv Preprint arXiv:1502.02127*.
- Dietterich, Thomas G. 2000. “Ensemble Methods in Machine Learning.” In *International Workshop on Multiple Classifier Systems*, 1–15. Springer.
- Eppstein, David, Maarten Löffler, and Darren Strash. 2010. “Listing All Maximal Cliques in Sparse Graphs in Near-Optimal Time.” In *International Symposium on Algorithms and Computation*, 403–14. Springer.
- Gelman, Andrew, Ben Goodrich, Jonah Gabry, and Aki Vehtari. 2019. “R-Squared for Bayesian Regression Models.” *The American Statistician*.
- Kapelner, Adam, and Justin Bleich. 2013. “bartMachine: Machine Learning with Bayesian Additive Regression Trees.” *arXiv Preprint arXiv:1312.2171*.
- . 2015. “Prediction with Missing Data via Bayesian Additive Regression Trees.” *Canadian Journal of Statistics* 43 (2): 224–39.
- Lipscomb, Carolyn E. 2000. “Medical Subject Headings (MeSH).” *Bulletin of the Medical Library Association* 88 (3): 265.
- Marshall, Iain J, and Byron C Wallace. 2019. “Toward Systematic Review Automation: A Practical Guide to Using Machine Learning Tools in Research Synthesis.” In *Systematic Reviews*, 8:1–10. 1. Springer.
- Miwa, Makoto, James Thomas, Alison O’Mara-Eves, and Sophia Ananiadou. 2014. “Reducing Systematic Review Workload Through Certainty-Based Screening.” *Journal of Biomedical Informatics* 51: 242–53.

- 212 “NCBI Insights : Updated Pubmed e-Utilities Coming in April 2022!” n.d. U.S. National Library of Medicine.
213 <https://ncbiinsights.ncbi.nlm.nih.gov/2021/10/05/updated-pubmed-api/>.
- 214 R Core Team. 2020. *R: A Language and Environment for Statistical Computing*. Vienna, Austria: R
215 Foundation for Statistical Computing. <https://www.R-project.org/>.
- 216 Robert, Christian P, George Casella, and George Casella. 2004. *Monte Carlo Statistical Methods*. Vol. 2.
217 Springer.
- 218 Rousseau, Francois. 2015. “Graph-of-Words: Mining and Retrieving Text with Networks of Features.” PhD
219 thesis, Ph. D. dissertation.
- 220 Settles, Burr. 2009. “Active Learning Literature Survey.”
- 221 Therneau, Terry, and Beth Atkinson. 2019. *Rpart: Recursive Partitioning and Regression Trees*. <https://CRAN.R-project.org/package=rpart>.
222
- 223 Wallace, Byron C, Kevin Small, Carla E Brodley, and Thomas A Trikalinos. 2010. “Active Learning for
224 Biomedical Citation Screening.” In *Proceedings of the 16th ACM SIGKDD International Conference on*
225 *Knowledge Discovery and Data Mining*, 173–82.
- 226 Yang, Li, and Abdallah Shami. 2020. “On Hyperparameter Optimization of Machine Learning Algorithms:
227 Theory and Practice.” *Neurocomputing* 415: 295–316.
- 228 Zhou, Zhi-Hua. 2021. “Ensemble Learning.” In *Machine Learning*, 181–210. Springer.